

## Urvashi Khandelwal

---

CONTACT INFORMATION	353 Jane Stanford Way Gates 232 Stanford CA, 94305	<a href="mailto:urvashik@stanford.edu">urvashik@stanford.edu</a> <a href="https://nlp.stanford.edu/~urvashik/">https://nlp.stanford.edu/~urvashik/</a>
EDUCATION	<b>Stanford University, CA</b> Ph.D. Computer Science Advisor: Dan Jurafsky	2015 - Present
	<b>University of Illinois Urbana-Champaign, IL</b> B.S. Computer Science, Minor in Mathematics Advisor: Jiawei Han	2011 - 2015
RESEARCH INTERESTS	Understanding and improving the generalization capabilities of neural language models and models for sequence generation; Non-parametric methods for improving out-of-distribution performance; Model interpretability and algorithmic fairness; Representation learning for better language understanding.	
AWARDS	<b>Microsoft Research Dissertation Grant</b> Winner	2020
	<b>CRA Outstanding Undergraduate Researchers Award</b> National Winner	2015
	<b>C.W. Gear Outstanding Undergraduate Award</b> Winner, University of Illinois Urbana-Champaign	2015
PUBLICATIONS	Nearest Neighbor Machine Translation. <b>Urvashi Khandelwal</b> , Angela Fan, Dan Jurafsky, Luke Zettlemoyer and Mike Lewis. International Conference on Learning Representations (ICLR), 2021.	
	With Little Power Comes Great Responsibility. Dallas Card, Peter Henderson, <b>Urvashi Khandelwal</b> , Robin Jia, Kyle Mahowald and Dan Jurafsky. Empirical Methods in Natural Language Processing (EMNLP), 2020.	
	Emergent Linguistic Structure in Artificial Neural Networks Trained by Self-Supervision. Chris Manning, Kevin Clark, John Hewitt, <b>Urvashi Khandelwal</b> and Omer Levy. Proceedings of the National Academy of Sciences (PNAS), 2020.	
	Generalization through Memorization: Nearest Neighbor Language Models. <b>Urvashi Khandelwal</b> , Omer Levy, Dan Jurafsky, Luke Zettlemoyer and Mike Lewis. International Conference on Learning Representations (ICLR), 2020.	
	What does BERT look at? An Analysis of BERT's Attention. Kevin Clark, <b>Urvashi Khandelwal</b> , Omer Levy and Christopher D. Manning. BlackboxNLP, 2019. ( <b>Best Paper Award</b> )	
	BAM! Born-Again Multi-Task Networks for Natural Language Understanding. Kevin Clark, Minh-Thang Luong, <b>Urvashi Khandelwal</b> , Christopher D. Manning	

and Quoc V. Le.  
Association for Computational Linguistics (ACL), 2019.

Sample Efficient Text Summarization Using a Single Pre-Trained Transformer.  
**Urvashi Khandelwal**, Kevin Clark, Dan Jurafsky, Lukasz Kaiser.  
ArXiv Preprint, 2019. Presented at WestCoast NLP, 2019.

Sharp Nearby, Fuzzy Far Away: How Neural Language Models Use Context.  
**Urvashi Khandelwal**, He He, Peng Qi and Dan Jurafsky.  
Association for Computational Linguistics (ACL), 2018.

ClusCite: Effective Citation Recommendation by Information Network Based Clustering.  
Xiang Ren, Jialu Liu, Xiao Yu, **Urvashi Khandelwal**, Quanquan Gu, Lidan Wang and Jiawei Han.  
International Conference on Knowledge Discovery and Data Mining (KDD), 2014.

Personalized Entity Recommendation in Heterogeneous Information Networks with Implicit User Feedback.  
Xiao Yu, Xiang Ren, Yizhou Sun, Quanquan Gu, Bradley Sturt, **Urvashi Khandelwal**, Brandon Norick and Jiawei Han.  
International Conference on Web Search and Data Mining, (WSDM), 2014.

HeteRec: Entity Recommendation in Heterogeneous Information Networks with Implicit User Feedback.  
Xiao Yu, Xiang Ren, Yizhou Sun, Bradley Sturt, **Urvashi Khandelwal**, Quanquan Gu, Brandon Norick, and Jiawei Han.  
International Conference on Recommender Systems (RecSys), 2013.

#### REPORTS

Government by Algorithm: Artificial Intelligence in Federal Administrative Agencies.  
Report submitted to the Administrative Conference of the United States (ACUS).  
February 2020.  
Served as the technical lead for two case studies:

- Informal Adjudication at the U.S. Patent and Trademark Office  
Daniel E. Ho, **Urvashi Khandelwal**, Alex Yu
- Formal Adjudication at the Social Security Administration  
Daniel E. Ho, Derin McLeod, **Urvashi Khandelwal**, Liza Starr, Emma Wang

#### INVITED TALKS

**Generalization through Memorization: Nearest Neighbor Language Models**  
Berkeley NLP Seminar. November, 2019.

**Sharp Nearby, Fuzzy Far Away: How Neural Language Models Use Context**  
Bay Area Research in NLP and ML Meetup. March, 2019.

**Media Portrayals of AI**  
Stanford AI Lab - AI Salon. April, 2017.

**Neural Text Summarization**  
Stanford Data Science Initiative. October, 2016.

#### PRESS COVERAGE

Facebook’s AI speeds up natural language processing without additional training.  
VentureBeat, February 19, 2020.

Helpful Neighbors.  
The Batch by Andrew Ng. January 29, 2019.

Stanford policy lab explores government use of artificial intelligence.  
Stanford News Service, February 28, 2019.

Khandelwal receives CRA Outstanding Undergraduate Researcher award.  
Department of Computer Science - CS@Illinois News, January 23, 2015.

## SKILLS

Languages: Python, C++, C, MATLAB, L<sup>A</sup>T<sub>E</sub>X  
Frameworks: PyTorch, Tensorflow

## WORK EXPERIENCE

### Facebook

Research Intern - Facebook AI Research Summer, 2020  
Managers: Luke Zettlemoyer, Mike Lewis

Worked on  $k$ -nearest neighbor machine translation models ( $k$ NN-MT), an extension of nearest neighbor language models to conditional generation tasks. Through a series of carefully designed experiments, we show that retrieving nearest neighbors from a datastore containing translation examples can improve model generalization without additional training or data. It can also be used to adapt translation systems to new domains by adding domain-specific data to the datastore, and to specialize multilingual translation systems to specific language pairs by adding language-specific data to the datastore, all without any added training costs.

### Facebook

Research Intern - Facebook AI Research Summer, 2019  
Managers: Mike Lewis, Omer Levy

Worked on  $k$ -nearest neighbor language models ( $k$ NN-LM) which involved extending pre-trained autoregressive LMs with a module to retrieve the nearest examples, in the pre-trained LM's representation space, from any text collection without any further training. This technique allows the model to generalize better using not only the LM's training data, but also by scaling up to larger corpora and adapting to different domains, by simply varying the text collection, without any additional training.

### Google

Research Intern - Google Brain Summer-Fall, 2018  
Manager: Lukasz Kaiser

Worked on large pre-trained language models in the context of language generation and multi-task learning. Demonstrated the sample efficiency of a pre-trained decoder-only Transformer for abstractive text summarization by concatenating the source article and the target summary and using the decoder to process both. The model achieved surprising performance from fine-tuning on only 3,000 training examples. In addition, collaborated on born-again multi-task learning for fine-tuning BERT on a number of language understanding tasks (the GLUE benchmark). A single multi-task model is fine-tuned via knowledge distillation from several single task models. This method consistently improves over standard single-task and multi-task training.

### Facebook

Software Engineering Intern - Data and Targeting, Ads Summer, 2014

Designed, implemented and deployed an A/B testing framework for clustering similar users with the goal of ads targeting. The aim of the project was to quantify, analyze and track gains and losses introduced by the development and deployment of new clus-

ters by different teams within the organization. The project included two main components: a recurring automated data extraction, processing and metrics calculation pipeline implemented in Hive and a front end interface for launching and monitoring A/B experiments and compiling the results into automatically generated visualizations. The correctness and utility of the project was verified by studying the impact of the pipeline on real users in live experiments, which led to significant increases in revenue for the specific categories that made use of them.

### Google

Software Engineering Intern - Gmail Backend Summer, 2013

Developed a performance analysis framework for the Gmail backend pipeline. The framework involved the design, implementation and testing of an anonymizer module as well as its integration with the existing Bigtable setup, where user privacy was a key concern. The module, implemented in C++, was designed to meet timing and memory constraints.

### Qualcomm

Engineering Intern - Qualcomm CDMA Technologies Summer, 2012

Designed Production Level APIs for the Camera Sensor Driver team to debug Hardware Software Interface issues. The APIs penetrated the high level Software in C++ to the lower Hardware layer code in C, using registers and I2C Reads to retrieve the information from the Sensor.

## TEACHING EXPERIENCE

### CS124 - From Languages to Information

Head Teaching Assistant, Stanford Winter 2019, 2020

### CS225 - Data Structures

Teaching Assistant, UIUC Spring 2013 - Fall 2014

### CS173 - Discrete Mathematics

Teaching Assistant, UIUC Fall 2013

### ECE110 - Introduction to Electrical and Computer Engineering

Teaching Assistant, UIUC Spring 2012

## SERVICE

### NeuralGen 2019 - Workshop on Methods for Optimizing and Evaluating Neural Language Generation

Co-Organizer 2019

### Stanford Computer Science PhD Admissions

Committee Member 2018-2019

### SAIL Undergraduate Mentorship Program

Mentor 2018-2019

### Stanford Women in AI

Organizer Fall 2017

### Stanford AI4ALL

NLP Research Project Mentor Summer, 2016

### Reviewer

ACL, NAACL, EMNLP, NeuralGen (Meta-reviewer), DeepGen, ACL-Student Research Workshop, KDD, RecSys